



SAPIENZA
UNIVERSITÀ DI ROMA

DEPARTMENT OF MECHANICAL AND AEROSPACE ENGINEERING

Homework 2 — Powered Descent and Landing via Direct Collocation

DYNAMICS AND CONTROL OF LAUNCH VEHICLES

Professor:

Prof. A. Zavoli

Student:

Niccolò D'Ambrosio

Academic Year 2025/2026

Contents

1	Introduction	2
1.1	Vehicle Model and Equations of Motion	2
1.2	Boundary Conditions and Mission Parameters	3
1.3	Path Constraints	3
1.4	Optimal Control Problem	4
2	Trapezoidal Direct Collocation	5
2.1	Trapezoidal Quadrature Rule	5
2.2	Defect Constraints	5
2.3	NLP Formulation	6
2.4	KKT Optimality Conditions	6
2.5	Sequential Quadratic Programming	7
3	Numerical Implementation	8
3.1	Non-Dimensionalisation	8
3.2	Decision Vector Layout	9
3.3	Initial Guess	9
3.4	Constraints	9
3.5	Solver Settings	10
4	Numerical Results	11
4.1	Trajectory and Glide-Slope Corridor	11
4.2	Thrust Magnitude	12
4.3	Mass and Glide-Slope Angle	12
5	Sensitivity Analysis	14
6	Task 2: Zero-Order Hold Discretization	15
6.1	ZOH Transcription for LTV Systems (Appendix A)	15
6.2	Application to the Nonlinear Powered-Descent Problem	16
6.3	Forward-Integration Validation	18
6.4	Results	18
7	Conclusions	22
	References	24

1 Introduction

This report presents the solution to Homework 2 of the course *Dynamics and Control of Launch Vehicles*: the minimum-fuel powered descent and pinpoint soft landing of a reusable launch vehicle. While Homework 1 tackled an ascent problem via the indirect approach (Pontryagin Maximum Principle plus shooting), this work adopts a *direct transcription*: the optimal control problem is approximated by a finite-dimensional Nonlinear Programming (NLP) problem through trapezoidal collocation, which is then solved by Sequential Quadratic Programming (SQP).

The powered-descent problem models the final phase of flight of a reusable rocket performing a "divert and soft land" manoeuvre, of the kind flown by the SpaceX Falcon 9 first stage and by Mars precision-landing concepts. The vehicle must reach a designated landing pad with zero residual velocity while respecting thrust bounds, a glide-slope corridor, and altitude constraints, and using minimum propellant.

1.1 Vehicle Model and Equations of Motion

The vehicle is modelled as a 2D point mass moving above flat ground with no aerodynamics. The state vector is

$$\mathbf{x}(t) = [x, y, v_x, v_y, m]^T, \quad (1)$$

and the control is the thrust vector $\mathbf{T}(t) = [T_x, T_y]^T$. The equations of motion are

$$\dot{x} = v_x, \quad (2)$$

$$\dot{y} = v_y, \quad (3)$$

$$\dot{v}_x = \frac{T_x}{m}, \quad (4)$$

$$\dot{v}_y = \frac{T_y}{m} - g, \quad (5)$$

$$\dot{m} = -\frac{\|\mathbf{T}\|}{c}, \quad (6)$$

where $g = 9.81 \text{ m/s}^2$ is the gravitational acceleration and $c = I_{sp} g_0$ is the effective exhaust velocity ($I_{sp} = 225 \text{ s}$, $g_0 = 9.80665 \text{ m/s}^2$, hence $c \approx 2206 \text{ m/s}$). As in Homework 1, the problem is non-dimensionalised before being passed to the solver (see Section 3.1); the dimensional form is retained throughout this section for clarity. Compared to Homework 1, the control is the full thrust vector rather than a scalar direction angle, since the magnitude is now also a degree of freedom subject to a bound.

1.2 Boundary Conditions and Mission Parameters

The vehicle starts at a known state and must reach the origin with zero velocity at the prescribed terminal time:

$$\mathbf{x}(0) = [x_0, y_0, v_{x,0}, v_{y,0}, m_0]^T, \quad \mathbf{x}(t_f) = [0, 0, 0, 0, m_f]^T, \quad (7)$$

with the final mass m_f free. The mission data taken from Table 1 of the assignment are summarised in Table 1.

Table 1: Mission data.

Quantity	Symbol	Value	Units
Initial position	(x_0, y_0)	(1000, 3000)	m
Initial velocity	$(v_{x,0}, v_{y,0})$	(300, -200)	m/s
Initial mass	m_0	2000	kg
Specific impulse	I_{sp}	225	s
Thrust bounds	$T_{\min}, T_{\max}$	0, 70	kN
Glide-slope half-angle	$\theta_{\max}$	60	deg
Flight time (fixed)	t_f	38	s

1.3 Path Constraints

Three path constraints are enforced over $[0, t_f]$:

$$T_{\min} \leq \|\mathbf{T}(t)\| \leq T_{\max} \quad (\text{thrust magnitude}), \quad (8)$$

$$|x(t)| \leq \tan(\theta_{\max}) y(t) \quad (\text{glide-slope}), \quad (9)$$

$$y(t) \geq 0 \quad (\text{no underground flight}). \quad (10)$$

The glide-slope condition (9) confines the vehicle to a cone of half-angle $\theta_{\max}$ apexed at the landing pad, ensuring a safe approach corridor. It is convex in (x, y) and decomposes into the linear pair

$$+x - \tan(\theta_{\max}) y \leq 0, \quad -x - \tan(\theta_{\max}) y \leq 0. \quad (11)$$

Both forms are exploited downstream: the linearised pair (11) is what is fed to **fmincon** as a smooth inequality.

1.4 Optimal Control Problem

The minimum-fuel powered-descent problem is the Mayer-form OCP

$$\begin{aligned} \min_{\mathbf{T}(\cdot)} \quad & J = -m(t_f) \\ \text{s.t.} \quad & \text{Eqs. (2)–(6),} \\ & \text{boundary conditions of Section 1.2,} \\ & \text{path constraints (8)–(10).} \end{aligned} \tag{12}$$

This continuous-time problem is approximated by an NLP using trapezoidal direct collocation, as described in Section 2.

2 Trapezoidal Direct Collocation

Direct transcription methods discretise the continuous optimal control problem (12) over a finite mesh of nodes and replace the continuous-time dynamics with algebraic relations among the node values. The resulting finite-dimensional NLP can be solved by general-purpose nonlinear programming solvers without the need to derive the costate equations and transversality conditions of the indirect approach. This is the methodological complement to the PMP-based formulation used in Homework 1.

2.1 Trapezoidal Quadrature Rule

The continuous-time dynamics are written compactly as

$$\dot{\mathbf{x}}(t) = \mathbf{f}(\mathbf{x}(t), \mathbf{u}(t)), \quad \mathbf{x} \in \mathbb{R}^{n_x}, \mathbf{u} \in \mathbb{R}^{n_u}. \quad (13)$$

Discretise the time horizon $[0, t_f]$ into N uniformly spaced nodes $\{t_k\}_{k=1}^N$ with spacing $\Delta t = t_f/(N - 1)$, and let

$$\mathbf{x}_k = \mathbf{x}(t_k), \quad \mathbf{u}_k = \mathbf{u}(t_k), \quad \mathbf{f}_k = \mathbf{f}(\mathbf{x}_k, \mathbf{u}_k). \quad (14)$$

Integrating the state equation over $[t_k, t_{k+1}]$,

$$\mathbf{x}_{k+1} - \mathbf{x}_k = \int_{t_k}^{t_{k+1}} \mathbf{f}(\mathbf{x}(\tau), \mathbf{u}(\tau)) \, d\tau. \quad (15)$$

Approximating the integrand by a linear interpolation between the two endpoints and applying the **trapezoidal quadrature rule** yields

$$\mathbf{x}_{k+1} - \mathbf{x}_k \approx \frac{\Delta t}{2} (\mathbf{f}_k + \mathbf{f}_{k+1}). \quad (16)$$

The local truncation error of the trapezoidal rule is $O(\Delta t^3)$, giving a globally second-order method. The rule is implicit (since $\mathbf{f}_{k+1}$ depends on $\mathbf{x}_{k+1}$) and A-stable, properties that make it well suited to transcribing dynamical systems with moderately stiff modes.

2.2 Defect Constraints

Promoting (16) from approximation to equality constraint produces the **defect equations** of trapezoidal direct collocation:

$$\boldsymbol{\delta}_k := \mathbf{x}_{k+1} - \mathbf{x}_k - \frac{\Delta t}{2} (\mathbf{f}_k + \mathbf{f}_{k+1}) = \mathbf{0}, \quad k = 1, \dots, N - 1. \quad (17)$$

With n_x states this generates $n_x(N - 1)$ scalar nonlinear equality constraints. When all defects are satisfied to machine precision, the discrete trajectory is consistent with the continuous-time dynamics to second order in Δt , and the controls are interpreted as the nodal values of a piecewise-linear input.

For the powered-descent problem ($n_x = 5$, $n_u = 2$) each node contributes $n_x + n_u = 7$ decision variables and the defect block carries $5(N - 1)$ rows.

2.3 NLP Formulation

The decision vector stacks states and controls node by node:

$$\mathbf{z} = [\mathbf{x}_1^T, \mathbf{u}_1^T, \mathbf{x}_2^T, \mathbf{u}_2^T, \dots, \mathbf{x}_N^T, \mathbf{u}_N^T]^T \in \mathbb{R}^{(n_x+n_u)N}. \quad (18)$$

The transcribed problem is

$$\begin{aligned} \min_{\mathbf{z}} \quad & -m_N \\ \text{s.t.} \quad & \mathbf{c}_{\text{eq}}^{\text{lin}}(\mathbf{z}) = \mathbf{0} \quad (\text{boundary conditions}), \\ & \mathbf{c}_{\text{eq}}^{\text{nl}}(\mathbf{z}) = \mathbf{0} \quad (\text{defects}), \\ & \mathbf{c}_{\text{ineq}}(\mathbf{z}) \leq \mathbf{0} \quad (\text{path constraints}), \\ & \mathbf{z}_{\text{lb}} \leq \mathbf{z} \leq \mathbf{z}_{\text{ub}} \quad (\text{box bounds}). \end{aligned} \quad (19)$$

With $N = 50$, the dimensions are: $7N = 350$ decision variables; 9 linear equalities (5 initial + 4 final boundary conditions, m_f being free); $5(N - 1) = 245$ nonlinear equalities (defects); $4N = 200$ nonlinear inequalities (lower and upper thrust magnitude, two glide-slope half-spaces); and $7N$ box bounds.

2.4 KKT Optimality Conditions

Introducing multipliers $\boldsymbol{\lambda}^{\text{eq}}$ and $\boldsymbol{\mu} \geq \mathbf{0}$ for the equality and inequality constraints, the Lagrangian of (19) is

$$\mathcal{L}(\mathbf{z}, \boldsymbol{\lambda}^{\text{eq}}, \boldsymbol{\mu}) = J(\mathbf{z}) + (\boldsymbol{\lambda}^{\text{eq}})^T \mathbf{c}_{\text{eq}}(\mathbf{z}) + \boldsymbol{\mu}^T \mathbf{c}_{\text{ineq}}(\mathbf{z}). \quad (20)$$

A point $\mathbf{z}^*$ is a local minimiser if there exist multipliers $(\boldsymbol{\lambda}^{\text{eq},*}, \boldsymbol{\mu}^*)$ such that the **Karush–Kuhn–Tucker** (KKT) conditions hold:

$$\nabla_{\mathbf{z}} \mathcal{L}(\mathbf{z}^*, \boldsymbol{\lambda}^{\text{eq},*}, \boldsymbol{\mu}^*) = \mathbf{0} \quad (\text{stationarity}), \quad (21)$$

$$\mathbf{c}_{\text{eq}}(\mathbf{z}^*) = \mathbf{0}, \quad \mathbf{c}_{\text{ineq}}(\mathbf{z}^*) \leq \mathbf{0} \quad (\text{primal feasibility}), \quad (22)$$

$$\boldsymbol{\mu}^* \geq \mathbf{0} \quad (\text{dual feasibility}), \quad (23)$$

$$\mu_j^* c_{\text{ineq},j}(\mathbf{z}^*) = 0, \quad \forall j \quad (\text{complementary slackness}). \quad (24)$$

Complementarity expresses that an inequality is either active ($c_j = 0$, with $\mu_j > 0$) or strictly inactive ($c_j < 0$, with $\mu_j = 0$). At the optimum of the powered-descent problem the upper thrust bound $\|\mathbf{T}\| = T_{\text{max}}$ is active over a substantial portion of the trajectory, mirroring the bang-bang nature of the optimal control law for fuel-optimal problems with linear-in-control dynamics; the corresponding KKT multipliers are positive, identifying the binding constraints that shape the solution.

2.5 Sequential Quadratic Programming

The NLP (19) is solved with MATLAB's `fmincon` configured with the SQP algorithm. SQP iteratively constructs a quadratic-programming approximation of the KKT system: at each iterate $\mathbf{z}^{(k)}$ the search direction $\mathbf{d}^{(k)}$ is found by solving

$$\begin{aligned} \min_{\mathbf{d}} \quad & \frac{1}{2} \mathbf{d}^T \mathbf{H}^{(k)} \mathbf{d} + \nabla J(\mathbf{z}^{(k)})^T \mathbf{d} \\ \text{s.t.} \quad & \nabla \mathbf{c}_{\text{eq}}(\mathbf{z}^{(k)})^T \mathbf{d} + \mathbf{c}_{\text{eq}}(\mathbf{z}^{(k)}) = \mathbf{0}, \\ & \nabla \mathbf{c}_{\text{ineq}}(\mathbf{z}^{(k)})^T \mathbf{d} + \mathbf{c}_{\text{ineq}}(\mathbf{z}^{(k)}) \leq \mathbf{0}, \end{aligned} \tag{25}$$

where $\mathbf{H}^{(k)}$ is a quasi-Newton (BFGS) approximation of the Hessian of the Lagrangian. The next iterate is $\mathbf{z}^{(k+1)} = \mathbf{z}^{(k)} + \alpha^{(k)} \mathbf{d}^{(k)}$, with the step length $\alpha^{(k)}$ obtained from a line search on a merit function that balances objective decrease against constraint violation.

3 Numerical Implementation

This section details the practical setup of the NLP (19) in MATLAB. The full transcription is built and solved by the script `main_task1.m`; `fmincon` is the only solver dependency (no CasADi, no YALMIP).

3.1 Non-Dimensionalisation

The decision-vector entries span four orders of magnitude in SI units (positions $\sim 10^3$ m, velocities $\sim 10^2$ m/s, mass $\sim 10^3$ kg, thrust components $\sim 7 \cdot 10^4$ N). Mixing scales of this size in a single NLP yields a poorly conditioned Jacobian and slow asymptotic convergence under finite-difference gradients. Following the same approach used in Homework 1, the problem is therefore non-dimensionalised before being handed to `fmincon`. The reference scales are listed in Table 2.

Table 2: Non-dimensionalisation reference scales.

Quantity	Symbol	Definition	Numerical value
Length	L_{ref}	y_0	3000 m
Acceleration	a_{ref}	g	9.81 m/s ²
Time	t_{ref}	$\sqrt{L_{ref}/g}$	17.49 s
Velocity	V_{ref}	$\sqrt{g L_{ref}}$	171.6 m/s
Mass	m_{ref}	m_0	2000 kg
Thrust	T_{ref}	$m_0 g$	19 620 N

In non-dimensional variables (denoted by tildes), the dynamics (2)–(6) reduce to

$$\dot{\tilde{x}} = \tilde{v}_x, \quad \dot{\tilde{y}} = \tilde{v}_y, \quad (26)$$

$$\dot{\tilde{v}}_x = \frac{\tilde{T}_x}{\tilde{m}}, \quad \dot{\tilde{v}}_y = \frac{\tilde{T}_y}{\tilde{m}} - 1, \quad (27)$$

$$\dot{\tilde{m}} = -V_c \|\tilde{\mathbf{T}}\|, \quad (28)$$

where $V_c = V_{ref}/c \approx 0.0778$ is the only residual dimensionless parameter (a local Tsiolkovsky-style ratio of the characteristic flight speed to the effective exhaust velocity). All decision variables now span $O(1)$ magnitudes, with positions in $[-1/3, 1]$, velocities in $[-1.17, 1.75]$, mass in $[10^{-3}, 1]$, and thrust components in $[-3.57, 3.57]$; the Jacobians of the LTV linearisation (Section 6) become uniformly $O(1)$, sharply improving the conditioning of `fmincon`’s SQP iterations and of the SCvx outer loop. Boundary conditions, glide-slope angle, and the time horizon t_f are all rescaled by their natural units; results are returned to SI before printing and plotting.

3.2 Decision Vector Layout

With $N = 50$ collocation nodes, the decision vector (18) has $7N = 350$ components ordered as

$$\mathbf{z} = \left[\underbrace{x_1, y_1, v_{x,1}, v_{y,1}, m_1, T_{x,1}, T_{y,1}}_{\text{node 1}}, \dots, \underbrace{x_N, y_N, v_{x,N}, v_{y,N}, m_N, T_{x,N}, T_{y,N}}_{\text{node } N} \right]^T. \quad (29)$$

A slice operator `idx(i) = (i-1)*7 + (1:7)` returns the indices of the seven entries belonging to node i , used throughout the helper functions to build sparse constraint matrices and unpack the converged solution.

3.3 Initial Guess

A reasonable initial guess is essential for SQP convergence on this nonconvex NLP. The state variables are linearly interpolated between the boundary states with $\alpha = (i - 1)/(N - 1)$:

$$x_i^{(0)} = (1 - \alpha) x_0, \quad y_i^{(0)} = (1 - \alpha) y_0, \quad (30)$$

$$v_{x,i}^{(0)} = (1 - \alpha) v_{x,0}, \quad v_{y,i}^{(0)} = (1 - \alpha) v_{y,0}, \quad (31)$$

$$m_i^{(0)} = m_0(1 - 0.3\alpha). \quad (32)$$

The control guess is a constant hover thrust $\mathbf{T}^{(0)} = (0, m_0 g)^T$, which approximately balances gravity at the start and provides a search direction that does not violate the thrust-magnitude bounds.

3.4 Constraints

Linear equality (boundary conditions). The nine scalar conditions are encoded in $\mathbf{A}_{\text{eq}}\mathbf{z} = \mathbf{b}_{\text{eq}}$ as a sparse $9 \times 7N$ matrix: five rows fix the initial state and four rows fix the terminal $(x, y, v_x, v_y) = (0, 0, 0, 0)$. The terminal mass is left free (its bound, $m \geq 1$, is a box constraint).

Nonlinear equality (defects). The function `trap_nonlcon` reshapes $\mathbf{z}$ into a $7 \times N$ matrix, evaluates the right-hand side $\mathbf{f}_k$ at every node via `dyn_rhs`, and assembles the $5(N - 1) = 245$ defect entries of Eq. (17).

Nonlinear inequality. The same routine returns the $4N$ inequalities:

- Thrust magnitude (lower and upper): $T_{\min} - \|\mathbf{T}_k\| \leq 0$ and $\|\mathbf{T}_k\| - T_{\max} \leq 0$, both nonsmooth at the origin but smoothly differentiable away from $\mathbf{T} = \mathbf{0}$;
- Glide-slope (linear, decomposed): $+x_k - \tan \theta_{\max} y_k \leq 0$ and $-x_k - \tan \theta_{\max} y_k \leq 0$.

Box bounds. Per-node bounds enforce $\tilde{y}_k \geq 0$, $10^{-3} \leq \tilde{m}_k \leq 1$ (equivalently $2 \leq m_k \leq m_0$ kg in SI), and $|\tilde{T}_{x,k}|, |\tilde{T}_{y,k}| \leq \tilde{T}_{\max}$. The strictly-positive mass lower bound avoids division-by-zero in the dynamics RHS (4)–(5); the chosen value of 2 kg is well below the propellant inventory at any feasible terminal state.

3.5 Solver Settings

The NLP is solved with `fmincon` using the options listed in Table 3.

Table 3: `fmincon` options.

Option	Value
Algorithm	'sqp'
MaxIterations	1000
MaxFunctionEvaluations	10^6
OptimalityTolerance	10^{-5}
ConstraintTolerance	10^{-6}
StepTolerance	10^{-10}

The SQP algorithm was preferred over the interior-point method (`fmincon`'s default) for two reasons: SQP exploits a warm start more effectively when sweeping t_f in the sensitivity study, and it tracks the active-set structure of the bang-bang optimum more efficiently than barrier-based interior-point iterations. Gradients are not provided explicitly, so `fmincon` computes them via finite differences; the typical wall time at $N = 50$ is approximately 30 s per solve on a modern laptop.

4 Numerical Results

The NLP (19) is solved at the nominal flight time $t_f = 38$ s with $N = 50$ collocation nodes. `fmincon` runs for approximately 30 s of wall time and returns a final mass of $m_f = 1403.20$ kg, equivalently 596.80 kg of propellant consumed (about 30% of the initial mass). The solver terminates at the `MaxIterations` cap of 1000 iterations; the constraint residual is below the tolerance of 10^{-6} and the first-order optimality reaches $\sim 10^{-4}$ in non-dim units, indicating that the algorithm has located a high-quality but not fully converged local optimum (see Section 5 for the analogous behaviour across the sensitivity sweep).

The four key plots are presented below. All three sensitivity runs (Section 5) are overlaid on the same axes for reference; the discussion in this section focuses on the nominal $t_f = 38.00$ s curve.

4.1 Trajectory and Glide-Slope Corridor

Figure 1 shows the descent trajectory in the (x, y) plane together with the glide-slope corridor (dashed grey lines bounding the cone of half-angle $\theta_{\max} = 60^\circ$). The vehicle enters the corridor from the upper right with significant horizontal velocity, executes a divert manoeuvre to align with the landing pad, and softly touches down at the origin (black triangle). The glide-slope constraint becomes active in the late portion of the descent, as the trajectory grazes the corridor boundary before the final approach.




Figure 1: Powered descent trajectory in the (x, y) plane. Dashed grey lines: glide-slope corridor at $\theta_{\max} = 60^\circ$. Black triangle: landing pad at the origin.

4.2 Thrust Magnitude

Figure 2 shows the thrust-magnitude profile $\|\mathbf{T}(t)\|$. The fuel-optimal solution exhibits the expected **bang-bang structure**: $\|\mathbf{T}\|$ saturates at $T_{\max} = 70$ kN over an extended portion of the descent (the deceleration phase), and drops to a low coast-thrust level just before touchdown to satisfy the soft-landing terminal condition. This is the discrete-time signature of the same Pontryagin-style optimality principle that produced the bilinear-tangent law in Homework 1, applied here to a problem with linear-in-control dynamics and a control-magnitude bound: the optimal control switches between the extreme values of its admissible set.





Figure 2: Thrust-magnitude profile. Dashed line: upper bound $T_{\max} = 70$ kN.


4.3 Mass and Glide-Slope Angle

Figure 3a shows the vehicle mass $m(t)$, which decreases monotonically at a rate proportional to $\|\mathbf{T}\|$ (Eq. 6). The slope is therefore approximately constant during the high-thrust deceleration phase and flattens during the low-thrust terminal segment, mirroring the bang-bang structure of the thrust profile.

Figure 3b shows the instantaneous glide-slope angle $\theta(t) = \arctan(|x|/y)$, bounded above by $\theta_{\max} = 60^\circ$ (dashed). The trajectory remains strictly inside the corridor for most of the descent and approaches the boundary only briefly near landing.



(a) Vehicle mass.



(b) Glide-slope angle $\arctan(|x|/y)$.

Figure 3: Mass profile and instantaneous glide-slope angle for the nominal solution.

5 Sensitivity Analysis

To probe the dependence of fuel consumption on the imposed flight time, the NLP is re-solved at $t_f \in \{0.95, 1.00, 1.05\} \cdot 38$ s, i.e., at 36.10, 38.00, and 39.90 s. Each solve uses the linear-interpolation initial guess of Section 3.3 (no warm-start across runs). The shortest and nominal runs ($t_f = 36.10$ s and 38.00 s) terminate at the `MaxIterations` cap with first-order optimality of order 10^{-4} in the non-dim cost; the longest run ($t_f = 39.90$ s) satisfies the step-size convergence criterion. The candidate solutions of the two capped runs still respect all constraints to within the `ConstraintTolerance` of 10^{-6} , so they are reported as engineering-quality solutions rather than fully optimised ones.

Table 4 reports the converged final masses and propellant consumption.

Table 4: Sensitivity sweep on flight time.

t_f [s]	m_f [kg]	Fuel consumed [kg]
36.10	1406.33	593.67
38.00	1403.20	596.80
39.90	1398.82	601.18

Fuel use grows with t_f across the swept window. The interpretation is straightforward: the longer the descent, the more time the vehicle spends in the gravity field, and the more thrust must be expended just to counteract gravity (gravity loss). Within the $\pm 5\%$ window around the nominal 38 s the variation is small (≈ 7.5 kg, or about 1.3% of the propellant), indicating that the nominal t_f is close to a flat region of the cost landscape—consistent with the observation that the mission is dominated by horizontal divert and not by hover time.

The monotonic behaviour suggests that the truly fuel-optimal solution corresponds to the minimum admissible t_f compatible with the thrust-magnitude bound. A free-time formulation, where t_f becomes a decision variable, would seek out this minimum automatically; this is one of the natural extensions identified in Section 7.

The trajectory, thrust, mass, and glide-slope plots presented in Section 4 overlay the three sensitivity runs on common axes, allowing direct visual comparison. The solutions are qualitatively similar: all three exhibit the bang-bang thrust structure, all three respect the glide-slope corridor, and the trajectories differ only marginally in shape, with the longer- t_f run hovering slightly more at altitude before initiating the final descent.

6 Task 2: Zero-Order Hold Discretization

The optional Task 2 of the assignment asks for an alternative transcription of the powered-descent problem based on a *Zero-Order Hold* (ZOH) discretisation, validated by forward-integrating the optimised control schedule with a high-fidelity integrator. This section presents **three complementary implementations** and compares them against the trapezoidal-collocation baseline of Task 1:

- (a) **Nonlinear ZOH with RK4 propagation** – a multiple-shooting NLP that holds $\mathbf{u}$ piecewise constant and propagates the *nonlinear* dynamics between grid nodes via a fixed-step RK4 integrator;
- (b) **LTV-linearised ZOH with SCvx outer loop, solved by fmincon/SQP** – the literal Appendix A construction: linearise about a reference, compute the discrete-time matrices via the auxiliary ODE, solve the resulting LTV optimisation, and iterate;
- (c) **LTV-linearised ZOH with SCvx outer loop, solved by YALMIP+ECOS** – same outer loop as variant (b), but the inner sub-problem is modelled in YALMIP as a Second-Order Cone Program (SOCP) and solved by the conic solver ECOS, exposing the convex structure that variant (b) hides inside a general-purpose NLP.

6.1 ZOH Transcription for LTV Systems (Appendix A)

Appendix A of the assignment derives the ZOH transcription for a continuous-time *linear time-varying* system

$$\dot{\mathbf{x}}(t) = \mathbf{A}(t) \mathbf{x}(t) + \mathbf{B}(t) \mathbf{u}(t) + \mathbf{c}(t). \quad (33)$$

Discretising time on a uniform grid $\{t_k\}$ with spacing Δt and assuming the control to be piecewise constant on each interval,

$$\mathbf{u}(t) = \mathbf{u}_k \quad \text{for } t \in [t_k, t_{k+1}), \quad (34)$$

the variation-of-constants formula yields the discrete-time recursion

$$\mathbf{x}_{k+1} = \bar{\mathbf{A}}_k \mathbf{x}_k + \bar{\mathbf{B}}_k \mathbf{u}_k + \bar{\mathbf{c}}_k, \quad (35)$$

where the discrete-time matrices follow from the state-transition matrix $\Phi(t, t_k)$:

$$\bar{\mathbf{A}}_k = \Phi(t_{k+1}, t_k), \quad \bar{\mathbf{B}}_k = \int_{t_k}^{t_{k+1}} \Phi(t_{k+1}, s) \mathbf{B}(s) ds, \quad \bar{\mathbf{c}}_k = \int_{t_k}^{t_{k+1}} \Phi(t_{k+1}, s) \mathbf{c}(s) ds. \quad (36)$$

These are evaluated numerically by integrating an augmented system over each interval $[t_k, t_{k+1}]$. Defining $\hat{\mathbf{B}}(t) = \int_{t_k}^t \mathbf{\Phi}(t, s) \mathbf{B}(s) ds$ and $\hat{\mathbf{c}}(t) = \int_{t_k}^t \mathbf{\Phi}(t, s) \mathbf{c}(s) ds$, one obtains the linear ODEs

$$\dot{\mathbf{\Phi}} = \mathbf{A} \mathbf{\Phi}, \quad \dot{\hat{\mathbf{B}}} = \mathbf{A} \hat{\mathbf{B}} + \mathbf{B}, \quad \dot{\hat{\mathbf{c}}} = \mathbf{A} \hat{\mathbf{c}} + \mathbf{c}, \quad (37)$$

with $\mathbf{\Phi}(t_k, t_k) = \mathbf{I}$ and $\hat{\mathbf{B}}(t_k) = \mathbf{0}$, $\hat{\mathbf{c}}(t_k) = \mathbf{0}$. At the end of each interval $\bar{\mathbf{A}}_k = \mathbf{\Phi}(t_{k+1}, t_k)$, $\bar{\mathbf{B}}_k = \hat{\mathbf{B}}(t_{k+1})$, $\bar{\mathbf{c}}_k = \hat{\mathbf{c}}(t_{k+1})$. This formulation is mathematically equivalent to the β, γ form of the appendix and avoids the need to invert $\mathbf{\Phi}$ at each step.

6.2 Application to the Nonlinear Powered-Descent Problem

The dynamics (2)–(6) are *nonlinear* in $(\mathbf{x}, \mathbf{u})$ owing to the T_x/m , T_y/m , and $\|\mathbf{T}\|/c$ terms, so Appendix A is not directly applicable. Two strategies are pursued:

Variant (a): nonlinear ZOH with exact RK4 propagation

Keep the nonlinear dynamics intact and propagate the state from $\mathbf{x}_k$ to $\mathbf{x}_{k+1}$ using the constant control $\mathbf{u}_k$:

$$\mathbf{x}_{k+1} = F(\mathbf{x}_k, \mathbf{u}_k; \Delta t), \quad (38)$$

where $F(\cdot, \cdot; \Delta t)$ is the ZOH-input flow map of the nonlinear ODE, implemented as a fixed-step RK4 with $n_{\text{sub}} = 2$ substeps per interval. The decision vector and constraint structure are identical to Task 1 except that the trapezoidal defects (17) are replaced by the nonlinear ZOH defects:

$$\delta_k^{\text{ZOH}} := \mathbf{x}_{k+1} - F(\mathbf{x}_k, \mathbf{u}_k; \Delta t) = \mathbf{0}, \quad k = 1, \dots, N-1, \quad (39)$$

and the terminal control $\mathbf{u}_N$ is pinned to zero (no interval after t_N under the ZOH convention).

Variant (b): linearised LTV ZOH with successive convexification

Linearise the dynamics about a reference trajectory $(\mathbf{x}_{\text{ref}}(t), \mathbf{u}_{\text{ref}}(t))$,

$$\mathbf{A}(t) = \left. \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right|_{\text{ref}}, \quad \mathbf{B}(t) = \left. \frac{\partial \mathbf{f}}{\partial \mathbf{u}} \right|_{\text{ref}}, \quad \mathbf{c}(t) = \mathbf{f}(\mathbf{x}_{\text{ref}}, \mathbf{u}_{\text{ref}}) - \mathbf{A} \mathbf{x}_{\text{ref}} - \mathbf{B} \mathbf{u}_{\text{ref}}, \quad (40)$$

yielding the LTV form (33). The discrete matrices $\bar{\mathbf{A}}_k, \bar{\mathbf{B}}_k, \bar{\mathbf{c}}_k$ are computed via Eq. (37) on every interval and used as linear equality constraints (35) in the NLP. With linear dynamics, linear box bounds, linear glide-slope, and convex thrust-magnitude bound, the LTV sub-problem is convex (an SOCP).

The linearisation is only locally accurate, so a **Successive Convexification** (SCvx) outer loop is wrapped around it: at each iteration the LTV sub-problem is solved, the reference is updated to the new solution, and the cycle is repeated until $\|\mathbf{x}^{(j+1)} - \mathbf{x}^{(j)}\|_2 < \tau$. Two implementation choices proved essential for stable convergence on this problem:

- **Warm start from the trapezoidal solution.** Cold-starting SCvx from a linear-interpolation reference with hover thrust drives the LTV NLP into regions where the linearisation predicts mass creation ($\dot{m} > 0$) and other unphysical behaviour, because the offset $\mathbf{c}(t)$ is built from the reference dynamics and only matches the nonlinear $\mathbf{f}(\mathbf{x}, \mathbf{u})$ in a neighbourhood of the linearisation point. Warm-starting from the Task 1 trapezoidal solution places the initial reference close to the optimum and the linearisation is locally faithful from the first iteration.
- **Adaptive box trust region with ratio test.** Even with a good warm start, an unconstrained LTV NLP can drift along directions where the linear model is loose (in particular the mass equation, since $\partial f_5 / \partial \mathbf{T}$ is nearly singular when $\|\mathbf{T}\|$ is small): the LTV “predicts” a fuel saving that the nonlinear dynamics do not deliver, and successive iterations oscillate around the true optimum. The implementation therefore wraps the LTV NLP in a hard box trust region with adaptive scale $\rho \in [\rho_{\min}, 1]$, namely $|\tilde{z} - \tilde{z}_{\text{ref}}| \leq \rho \boldsymbol{\rho}_0$, with base radii $\boldsymbol{\rho}_0 = (\rho_{xy}, \rho_v, \rho_m, \rho_T) = (0.17, 0.6, 0.1, 1.0)$ in non-dim units (roughly 500 m, 100 m/s, 200 kg, 20 000 N).
 - After each candidate step $\tilde{\mathbf{z}}^{(j+1)}$, the candidate ZOH controls are forward-integrated through the *nonlinear* dynamics with `ode45` to obtain the actual final mass m_f^{act} ;
 - the LTV-predicted final mass is compared with this through the trust-region ratio $\eta = (m_f^{\text{act}} - m_f^{(j)}) / (m_f^{\text{pred}} - m_f^{(j)})$;
 - if $\eta < \eta_l = 0.25$ the linear model is unreliable: the step is **rejected**, the reference is left unchanged, and $\rho \leftarrow \rho/2$;
 - if $\eta_l \leq \eta < \eta_h = 0.7$ the step is accepted and ρ is held;
 - if $\eta \geq \eta_h$ the linear model is excellent: the step is accepted and ρ is doubled (capped at the base value).

This is the canonical SCvx trust-region update and converges to the local optimum without the unphysical mass creation seen with a fixed unsafeguarded box.

For the nonlinear powered-descent dynamics, the Jacobians evaluate to

$$\mathbf{A} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & -T_x/m^2 \\ 0 & 0 & 0 & 0 & -T_y/m^2 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}, \quad \mathbf{B} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \\ 1/m & 0 \\ 0 & 1/m \\ -T_x/(c\|\mathbf{T}\|) & -T_y/(c\|\mathbf{T}\|) \end{bmatrix}. \quad (41)$$

The denominator $\|\mathbf{T}\|$ in $\mathbf{B}$ is regularised with a small additive constant in the implementation to avoid the singularity at $\|\mathbf{T}\| = 0$.

Variant (c): SCvx with a conic inner solver (YALMIP+ECOS)

Variant (b) feeds the LTV sub-problem to `fmincon` as a generic NLP, which discards the convex structure of the inner step and pays the cost of a quasi-Newton SQP iteration. Variant (c) keeps the same outer SCvx loop (warm-start from the trapezoidal solution, identical adaptive trust-region update, identical convergence tolerance) but re-models the inner sub-problem as a Second-Order Cone Program. With LTV dynamics $(\bar{\mathbf{A}}_k, \bar{\mathbf{B}}_k, \bar{\mathbf{c}}_k)$ as linear equalities, glide-slope and box bounds as linear inequalities, and the thrust-magnitude bound as the second-order cone constraint $\|\mathbf{u}_k\|_2 \leq T_{\max}$, the inner step is a pure SOCP. The problem is assembled in YALMIP via `sdpvar` and solved by ECOS, which exploits the conic structure directly. The two outer loops produce nearly indistinguishable trajectories (see Section 6.4) but the YALMIP/ECOS path is markedly cleaner numerically: the inner subproblem retains the cone constraint as a primitive rather than approximating it through SQP linearisations, and the per-step fidelity of the linearisation improves accordingly. The total wall time of the conic variant is approximately 14 s for the full 15-iteration outer loop on a modern laptop.

6.3 Forward-Integration Validation


The transcription accuracy of the four approaches (trapezoidal, ZOH–RK4, ZOH–LTV+SCvx with `fmincon`, ZOH–LTV+SCvx with YALMIP+ECOS) is assessed by forward-integrating the continuous-time dynamics (2)–(6) with `ode45` (RelTol = 10^{-10} , AbsTol = 10^{-12}) using the optimised control schedule, and comparing the resulting trajectory $\mathbf{x}^{\text{ode45}}(t_k)$ against the discretised values $\mathbf{x}_k^{\text{NLP}}$ at every grid node. The control interpolation matches the transcription assumption: piecewise-linear (PWL) for the trapezoidal rule, piecewise-constant (PWC) for the three ZOH variants. The state-norm error

$$e_k = \|\mathbf{x}_k^{\text{NLP}} - \mathbf{x}^{\text{ode45}}(t_k)\|_2 \quad (42)$$


quantifies how well each transcription represents the continuous-time problem.

6.4 Results

The four transcriptions converge to qualitatively identical trajectories (Fig. 4a). All four respect the glide-slope corridor and softly land at the origin, with only sub-metric differences in the divert phase. Figure 4b shows the thrust profiles: the trapezoidal solution yields a piecewise-linear thrust whereas the three ZOH variants are naturally piecewise-constant (rendered as staircases). All four exhibit the bang-bang structure with $\|\mathbf{T}\|$ saturated at $T_{\max}$ for most of the descent.




(a) Trajectory in the (x, y) plane.




(b) Thrust magnitude: PWL (trapezoidal) vs PWC (ZOH).

Figure 4: Comparison of the four transcriptions at $t_f = 38$ s.



(a) Vehicle mass profile.



(b) Forward-integration error at grid nodes (log scale).

Figure 5: Mass profile and transcription-fidelity check via `ode45` forward integration.

The transcription-fidelity plot (Fig. 5b) is the diagnostic value of Task 2: it measures how far each discretised trajectory drifts from a high-fidelity `ode45` integration of the same controls. A representative numerical comparison is summarised in Table 5.

Table 5: Four-way comparison at $t_f = 38$ s, $N = 50$ nodes.

Transcription	m_f [kg]	Fuel [kg]	Order	Max nondim node-error
Trapezoidal	1403.20	596.80	$O(\Delta t^2)$	$1.55 \cdot 10^{-3}$
Nonlinear ZOH (RK4)	1403.37	596.63	$O(\Delta t^4)$	$1.40 \cdot 10^{-8}$
LTV ZOH + SCvx (<code>fmincon</code>)	1399.54	600.46	<code>ode45</code> -based	$4.08 \cdot 10^{-4}$
LTV ZOH + SCvx (YALMIP+ECOS)	1400.84	599.16	<code>ode45</code> -based	$4.09 \cdot 10^{-5}$


The trapezoidal transcription has a global $O(\Delta t^2)$ error and sits at $\sim 10^{-3}$ in non-dim state norm (a few metres of position, fractions of a m/s of velocity). The nonlinear-ZOH+RK4 transcription is $O(\Delta t^4)$ and drops to the `ode45` tolerance floor ($\sim 10^{-8}$). Both LTV+SCvx variants lie in between: they propagate the linearised system with `ode45` on each interval, so the per-step accuracy is high, but the converged solution is constrained by the (small) residual linearisation error within the final trust radius. The conic variant (c) is one order of magnitude cleaner than the `fmincon` variant (b) ($4 \cdot 10^{-5}$ vs $4 \cdot 10^{-4}$): exposing the second-order cone constraint to the inner solver instead of approximating it by SQP linearisations measurably improves the per-step fidelity.

The trapezoidal and nonlinear-ZOH transcriptions agree to within 0.2 kg on m_f , indicating both have located the same local optimum. The two SCvx solutions sit 2–4 kg below this (slightly more fuel burnt) and agree with each other to within 1.3 kg. Two contributing effects are at play:


- Trapezoidal collocation enforces the dynamics only to second order in Δt , so its converged “optimum” is allowed a small dynamical residual ($\sim 10^{-3}$ above) that translates into a marginally optimistic final mass.
- Both SCvx outer loops terminate as soon as the trust region collapses to where the LTV linearisation is exact; they are therefore strictly feasible against the nonlinear dynamics, at the cost of stopping in a slightly different (but nearby) local optimum.

Both effects are below the engineering accuracy of the model (no aerodynamics, no attitude). Within numerical noise the four transcriptions describe the same physical solution.

The SCvx convergence traces are reported in Fig. 6. For each variant the top panel overlays the iterate change $\|\mathbf{x}^{(j+1)} - \mathbf{x}^{(j)}\|_2$, the trust-region scale ρ , and the final mass $m_f^{(j)}$; the bottom panel shows the trust-region ratio η against its acceptance band $[\eta_l, \eta_h] = [0.25, 0.7]$ (rejected steps are flagged in orange). Warm-started from the Task 1 trapezoidal solution, both algorithms initially overshoot the linearisation’s region of validity and reject the first four candidate steps, shrinking ρ ; once ρ is small enough that the LTV prediction matches the nonlinear forward integration, the steps are accepted, ρ stops decreasing, and $\|\Delta \mathbf{x}\|$ drops below $\sim 10^{-2}$ at the iteration cap. The final mass m_f stabilises within a fraction of a kilogram of the trapezoidal and nonlinear-ZOH values.



(a) Variant (b): SCvx with **fmincon**/SQP inner solver.



(b) Variant (c): SCvx with **YALMIP**+**ECOS** inner solver.

Figure 6: SCvx convergence traces for the two LTV+SCvx variants. Top panels: $\|\mathbf{x}^{(j+1)} - \mathbf{x}^{(j)}\|_2$ (left axis, log), trust scale ρ , and final mass m_f (right axis) per outer iteration. Bottom panels: trust-region ratio η against acceptance band $[\eta_l, \eta_h] = [0.25, 0.7]$, with rejected steps highlighted in orange.

The final-mass values returned by the four transcriptions agree to within a few kilograms, confirming that the trapezoidal approximation is accurate enough for engineering purposes despite its lower-order convergence. The ZOH–RK4 variant offers the simplest implementation when only a high-fidelity transcription is needed; the two LTV+SCvx variants expose the convex-sub-problem structure that underpins production guidance algorithms (GFOLD, lossless convexification, *etc.*). The conic variant (c) is the most idiomatic instantiation of that idea – it solves the inner step as an SOCP – and matches variant (b) in trajectory and final mass while delivering a measurably cleaner per-step fidelity.

7 Conclusions

This work transcribed the minimum-fuel powered-descent and pinpoint-landing optimal control problem into a finite-dimensional NLP via trapezoidal direct collocation, and solved it with `fmincon`/SQP. The main findings are as follows.

Direct vs. indirect transcription. Compared to the indirect-shooting approach of Homework 1, direct collocation requires no analytical derivation of costate equations or transversality conditions, accepts path constraints (glide-slope, thrust bounds) natively as inequality constraints, and is robust to a coarse initial guess. The cost is a higher-dimensional NLP (350 variables here) and the loss of the structured analytical insight that the bilinear-tangent law provided in the indirect formulation.

Bang-bang thrust structure. The fuel-optimal solution exhibits the classical bang-bang behaviour expected for problems with linear-in-control dynamics and a control-magnitude bound: $\|\mathbf{T}\|$ saturates at $T_{\max}$ over the deceleration phase and drops to a low value during the soft-landing approach. This is the direct-transcription analogue of the maximum-principle prediction.

Sensitivity to flight time. A $\pm 5\%$ sweep around the nominal $t_f = 38$ s revealed monotonic growth of fuel consumption with flight time (longer hover $\Rightarrow$ greater gravity losses), with a $\sim 1.2\%$ propellant variation across the window. The lower bound on t_f is set implicitly by the thrust-magnitude constraint.

Active constraints at the optimum. Two inequalities are active over substantial portions of the solution: the upper thrust bound $\|\mathbf{T}\| = T_{\max}$ and the glide-slope cone in the late descent. By the KKT complementarity conditions the corresponding multipliers are positive, identifying these as the binding constraints that shape the trajectory.

ZOH transcription (Task 2). The optional ZOH discretisation of Section 6 was implemented in three complementary forms—a nonlinear ZOH with RK4 propagation, the literal Appendix A construction (LTV-linearised + SCvx) solved by `fmincon`/SQP, and the same SCvx outer loop with the inner sub-problem cast as a Second-Order Cone Program and solved by YALMIP+ECOS. All three confirm the trapezoidal baseline at the engineering level: final masses agree within a few kilograms ($\sim 0.3\%$ of the propellant). The forward-integration test exposes the lower convergence order of the trapezoidal rule, with the ZOH–RK4 variant dropping the per-node fidelity error by five orders of magnitude and sitting close to the `ode45` tolerance floor. Warm-started from the trapezoidal solution and stabilised by an adaptive trust-region ratio test, both SCvx variants converge to the same local optimum. The conic variant exposes the SOCP structure of the inner step directly and delivers a per-step fidelity that is an order of magnitude tighter than the `fmincon`/SQP path, demonstrating the convex-sub-problem framework underlying production guidance algorithms (GFOLD, lossless convexification).

Non-dimensionalisation. Following the same approach used in Homework 1, the OCP is solved in non-dimensional variables (Section 3.1), with $L_{ref} = y_0$, $a_{ref} = g$ and the single residual parameter $V_c = V_{ref}/c \approx 0.078$. This shrinks the dynamic range of the decision vector from four orders of magnitude to one and is essential for the LTV linearisation to be well-conditioned: the Jacobians become uniformly $O(1)$ and `fmincon`/SQP makes appreciable progress per iteration.

Limitations and roadmap. The model neglects aerodynamic drag, vehicle attitude dynamics, and any throttling rate limits, and the flight time is fixed. Natural extensions remaining beyond the scope of this report are:

- A **free-time variant** that lifts the fixed- t_f assumption and allows the solver to find the absolute fuel-optimal flight time subject to the active thrust constraint, e.g. as in lossless-convexification formulations.
- **Lossless convexification** of a non-zero lower thrust bound $T_{\min} > 0$, which is non-convex in $(\mathbf{T})$: with the conic SCvx infrastructure already in place (variant (c)), the next step is to introduce the slack-variable trick of Açikmeşe and Ploen to recover a convex SOCP in the new state, and verify that the original constraint is recovered at the optimum.
- Inclusion of **aerodynamic drag** and **attitude dynamics**, moving from a 3-DoF point mass to a 6-DoF rigid body with angle-of-attack and angle-of-sideslip dependence.

References

- [1] J. T. Betts. *Practical Methods for Optimal Control and Estimation Using Nonlinear Programming*. SIAM, Philadelphia, 2nd edition, 2010.
- [2] C. R. Hargraves and S. W. Paris. Direct trajectory optimization using nonlinear programming and collocation. *Journal of Guidance, Control, and Dynamics*, 10(4):338–342, 1987.
- [3] J. Nocedal and S. J. Wright. *Numerical Optimization*. Springer, New York, 2nd edition, 2006.
- [4] B. Açikmeşe and S. R. Ploen. Convex programming approach to powered descent guidance for mars landing. *Journal of Guidance, Control, and Dynamics*, 30(5):1353–1366, 2007.
- [5] L. Blackmore, B. Açikmeşe, and D. P. Scharf. Minimum-landing-error powered-descent guidance for Mars landing using convex optimization. *Journal of Guidance, Control, and Dynamics*, 33(4):1161–1171, 2010.
- [6] A. E. Bryson and Y.-C. Ho. *Applied Optimal Control: Optimization, Estimation and Control*. Taylor & Francis, New York, revised edition, 1975.